

VELES HACK



HYPER-AI



CoGNETs
Contributors of Smart AI

ENACT

OCEI

Hyperion Sentinel

The edge-readiness engineer inside the HYPER-AI IDE

Veles Hack 2026 | Challenge 1: Hyperion, an LLM-powered agentic assistant



Funded by
the European Union

GitHub repo

- **GitHub repo:** <https://github.com/TusharTechs/hyperion-sentinel>
- **Docker image:** <https://hub.docker.com/r/tushartechs/hyperion>
- **Demo video (2 min, recorded live in the real IDE):** <https://youtu.be/ZP7SHtDgncU>



Source code

github.com/TusharTechs/hyperion-sentinel

scan to open



Live demo

youtu.be/ZP7SHtDgncU

scan to open




Docker image

hub.docker.com/r/tushartechs/hyperion

scan to open

Summary

- **What:** an agentic assistant for the HYPER-AI IDE. You chat; it inspects your workspace and acts through IDE actions (create / edit / delete files).
- **Problem:** edge deployments fail on the basics: latest images, no limits or probes, root, secrets in config.
- **Solution:** a deterministic Edge Readiness analyzer finds and scores issues; the LLM only converses.
- **Loop:** inspect, plan, you approve, apply via the IDE, re-read: before vs after.
- **Criteria:** guardrails, RAG on the official docs, memory, human-in-the-loop.





Highlights

- **Verified live** in the real IDE with Llama 3.1: readiness 35 to 65, 12 findings fixed.
- **Deterministic**: the LLM never sets scores or findings; secrets are never printed.
- **Safe by design**: every delete, overwrite and plan needs an explicit yes; guardrails fail closed.
- **Grounded + memory**: RAG with sources; per-user memory ("fix the second issue").
- **Robust**: works without an IDE backend or API key; 200+ tests; non-root amd64 image.
- **Open source** (Apache-2.0): GitHub, Docker Hub, 2-minute live demo video.

35 → 65

edge readiness, verified live

12

findings resolved in one chat

0

changes without your explicit
yes

200+

automated tests, Docker
amd64